

assignment3

October 7, 2023

*You are currently looking at **version 0.1** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](#) course resource.*

```
[1]: import numpy as np
import pandas as pd
```

0.0.1 Question 1

Import the data from `assets/fraud_data.csv`. What percentage of the observations in the dataset are instances of fraud?

This function should return a float between 0 and 1.

```
[2]: def answer_one():
    # YOUR CODE HERE

    # Use X_train, X_test, y_train, y_test for all of the following questions
    from sklearn.model_selection import train_test_split

    df = pd.read_csv('assets/fraud_data.csv')

    X = df.iloc[:, :-1]
    y = df.iloc[:, -1]

    X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)

    return (sum(y_train) + sum(y_test)) / (len(y_train) + len(y_test))

    raise NotImplementedError()
```

```
[ ]:
```

```
[3]: answer_one()
```

```
[3]: 0.016410823768035772
```

```
[4]: # Use X_train, X_test, y_train, y_test for all of the following questions
from sklearn.model_selection import train_test_split

df = pd.read_csv('assets/fraud_data.csv')

X = df.iloc[:, :-1]
y = df.iloc[:, -1]

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

```
[5]: (sum(y_train) + sum(y_test))/(len(y_train) + len(y_test))
```

```
[5]: 0.016410823768035772
```

0.0.2 Question 2

Using `X_train`, `X_test`, `y_train`, and `y_test` (as defined above), train a dummy classifier that classifies everything as the majority class of the training data. What is the accuracy of this classifier? What is the recall?

This function should return a tuple with two floats, i.e. (accuracy score, recall score).

```
[6]: from sklearn.dummy import DummyClassifier
from sklearn.metrics import recall_score

dummy = DummyClassifier(strategy = 'most_frequent', random_state = 0)

dummy.fit(X_train, y_train)

score = dummy.score(X_test, y_test)

predictions = dummy.predict(X_test)
recall = recall_score(y_test, predictions)

print(score)
print(recall)
```

```
0.9852507374631269
0.0
```

```
[7]: def answer_two():
    from sklearn.dummy import DummyClassifier
    from sklearn.metrics import recall_score

    dummy = DummyClassifier(strategy = 'most_frequent', random_state = 0)

    dummy.fit(X_train, y_train)
```

```

score = dummy.score(X_test, y_test)

predictions = dummy.predict(X_test)
recall = recall_score(y_test, predictions)

return (score, recall)

# YOUR CODE HERE
raise NotImplementedError()

```

[]:

0.0.3 Question 3

Using `X_train`, `X_test`, `y_train`, `y_test` (as defined above), train a SVC classifier using the default parameters. What is the accuracy, recall, and precision of this classifier?

This function should return a tuple with three floats, i.e. (accuracy score, recall score, precision score).

```

[8]: from sklearn.metrics import recall_score, precision_score
      from sklearn.svm import SVC

      svc = SVC()

      svc.fit(X_train, y_train)

      accuracy = svc.score(X_test, y_test)

      recall = recall_score(y_test, svc.predict(X_test))

      precision = precision_score(y_test, svc.predict(X_test))

      print(accuracy, ', ', recall, ', ', precision)

```

0.9900442477876106 , 0.35 , 0.9333333333333333

```

[9]: def answer_three():
      from sklearn.metrics import recall_score, precision_score
      from sklearn.svm import SVC

      # YOUR CODE HERE

      from sklearn.metrics import recall_score, precision_score
      from sklearn.svm import SVC

      svc = SVC()

```

```

svc.fit(X_train, y_train)

accuracy = svc.score(X_test, y_test)

recall = recall_score(y_test, svc.predict(X_test))

precision = precision_score(y_test, svc.predict(X_test))

return (accuracy, recall, precision)

raise NotImplementedError()

```

[]:

[10]: answer_three()

[10]: (0.9900442477876106, 0.35, 0.9333333333333333)

0.0.4 Question 4

Using the SVC classifier with parameters {'C': 1e9, 'gamma': 1e-07}, what is the confusion matrix when using a threshold of -220 on the decision function. Use X_test and y_test.

This function should return a confusion matrix, a 2x2 numpy array with 4 integers.

```

[11]: from sklearn.metrics import confusion_matrix
      from sklearn.svm import SVC

      svc = SVC(C = 1e9, gamma = 1e-07)

      svc.fit(X_train, y_train)

      y_predictions = svc.decision_function(X_test)

      y_predictions = np.where(y_predictions > -220, 1, 0)

      confusion_matrix(y_test, y_predictions)

```

[11]: array([[5320, 24],
[14, 66]])

[12]: svc.decision_function(X_test)

[12]: array([-739.71796843, -1086.16794833, -696.46339735, ...,
-491.97916719, -699.03838333, -701.93409309])

```
[13]: def answer_four():
    from sklearn.metrics import confusion_matrix
    from sklearn.svm import SVC

    svc = SVC(C = 1e9, gamma = 1e-07)

    svc.fit(X_train, y_train)

    y_predictions = svc.decision_function(X_test)

    y_predictions = np.where(y_predictions > -220, 1, 0)

    return confusion_matrix(y_test, y_predictions)

    raise NotImplementedError()
```

```
[ ]:
```

```
[14]: answer_four()
```

```
[14]: array([[5320,  24],
           [ 14,  66]])
```

0.0.5 Question 5

Train a logistic regression classifier with default parameters using `X_train` and `y_train`.

For the logistic regression classifier, create a precision recall curve and a roc curve using `y_test` and the probability estimates for `X_test` (probability it is fraud).

Looking at the precision recall curve, what is the recall when the precision is 0.75?

Looking at the roc curve, what is the true positive rate when the false positive rate is 0.16?

This function should return a tuple with two floats, i.e. (recall, true positive rate).

```
[15]: from sklearn.linear_model import LogisticRegression
    from sklearn.metrics import precision_recall_curve
    from sklearn.metrics import roc_curve

    logreg = LogisticRegression()

    logreg.fit(X_train, y_train)


    precision, recall, threshold_pr = precision_recall_curve(y_test, logreg.
        ↪predict_proba(X_test)[: , 1])

    false_positive_rate, true_positive_rate, threshold_roc = roc_curve(y_test,
        ↪logreg.predict_proba(X_test)[: , 1])
```

```
[16]: import matplotlib.pyplot as plt

plt.fill_between(recall, precision)
plt.ylabel("Precision")
plt.xlabel("Recall")
plt.title("Precision-Recall curve")
```


```
[16]: Text(0.5, 1.0, 'Precision-Recall curve')
```



```
[17]: import matplotlib.pyplot as plt

plt.fill_between(false_positive_rate, true_positive_rate)
plt.ylabel("True Positive Rate")
plt.xlabel("False Positive Rate")
plt.title("ROC Curve")
```

```
[17]: Text(0.5, 1.0, 'ROC Curve')
```



```
[18]: def answer_five():

    # YOUR CODE HERE

    return (0.84, 0.86)

    raise NotImplementedError()
```

```
[ ]:
```

0.0.6 Question 6

Perform a grid search over the parameters listed below for a Logistic Regression classifier, using recall for scoring and the default 3-fold cross validation. (Suggest to use `solver='liblinear'`, more explanation [here](#))

```
'penalty': ['l1', 'l2']
```

```
'C': [0.01, 0.1, 1, 10]
```

From `.cv_results_`, create an array of the mean test scores of each parameter combination. i.e.

| | 11 | 12 |
|------|----|----|
| 0.01 | ? | ? |
| 0.1 | ? | ? |
| 1 | ? | ? |
| 10 | ? | ? |

This function should return a 4 by 2 numpy array with 8 floats.

Note: do not return a DataFrame, just the values denoted by ? in a numpy array.

```
[21]: def answer_six():
    from sklearn.model_selection import GridSearchCV
    from sklearn.linear_model import LogisticRegression

    # YOUR CODE HERE
    logreg = LogisticRegression().fit(X_train, y_train)

    grid_values = {'penalty': ['l1', 'l2'], 'C': [0.01, 0.1, 1, 10]}

    grid_classifier_recall = GridSearchCV(logreg, param_grid = grid_values,
    ↪scoring = 'recall')

    grid_classifier_recall.fit(X_train, y_train)

    return np.array([grid_classifier_recall.cv_results_['mean_test_score'][x:
    ↪x+2] for x in range(0, len(grid_classifier_recall.
    ↪cv_results_['mean_test_score']), 2)])

    raise NotImplementedError()
```

```
[ ]:
```

```
[22]: answer_six()
```

```
[22]: array([[ nan, 0.80064935],
             [ nan, 0.80428571],
             [ nan, 0.80422078],
             [ nan, 0.80792208]])
```

```
[24]: from sklearn.model_selection import GridSearchCV
    from sklearn.linear_model import LogisticRegression

    logreg = LogisticRegression().fit(X_train, y_train)

    grid_values = {'penalty': ['l1', 'l2'], 'C': [0.01, 0.1, 1, 10]}
```



```
grid_classifier_recall = GridSearchCV(logreg, param_grid = grid_values, scoring_
    ⇨='recall')

grid_classifier_recall.fit(X_train, y_train)
```

```
[24]: GridSearchCV(estimator=LogisticRegression(),
    param_grid={'C': [0.01, 0.1, 1, 10], 'penalty': ['l1', 'l2']},
    scoring='recall')
```

```
[25]: grid_classifier_recall.cv_results_
```

```
[25]: {'mean_fit_time': array([2.02025890e-02, 6.37615609e+00, 2.02903271e-02,
    6.71791191e+00,
    2.05406666e-02, 7.33892756e+00, 1.65262222e-03, 5.13848410e+00]),
    'std_fit_time': array([3.60238985e-02, 1.87589571e+00, 3.77959661e-02,
    1.85844788e+00,
    3.80328061e-02, 1.62697709e+00, 5.34945767e-04, 9.33233874e-01]),
    'mean_score_time': array([0.          , 0.08093534, 0.          , 0.12026129, 0.
    ,
    0.11920462, 0.          , 0.09956293]),
    'std_score_time': array([0.          , 0.03792358, 0.          , 0.03947502, 0.
    ,
    0.04014686, 0.          , 0.00049915]),
    'param_C': masked_array(data=[0.01, 0.01, 0.1, 0.1, 1, 1, 10, 10],
    mask=[False, False, False, False, False, False, False, False],
    fill_value='?',
    dtype=object),
    'param_penalty': masked_array(data=['l1', 'l2', 'l1', 'l2', 'l1', 'l2', 'l1',
    'l2'],
    mask=[False, False, False, False, False, False, False, False],
    fill_value='?',
    dtype=object),
    'params': [{'C': 0.01, 'penalty': 'l1'},
    {'C': 0.01, 'penalty': 'l2'},
    {'C': 0.1, 'penalty': 'l1'},
    {'C': 0.1, 'penalty': 'l2'},
    {'C': 1, 'penalty': 'l1'},
    {'C': 1, 'penalty': 'l2'},
    {'C': 10, 'penalty': 'l1'},
    {'C': 10, 'penalty': 'l2'}],
    'split0_test_score': array([
    nan, 0.78181818,
    nan, 0.78181818,
    0.78181818,
    nan, 0.78181818]),
    'split1_test_score': array([
    nan, 0.81818182,
    nan, 0.81818182,
    0.83636364,
    nan, 0.83636364]),
    'split2_test_score': array([
    nan, 0.87272727,
    nan, 0.89090909,
```

```

nan,
    0.89090909,      nan, 0.89090909]),
'split3_test_score': array([      nan, 0.82142857,      nan, 0.82142857,
nan,
    0.83928571,      nan, 0.82142857]),
'split4_test_score': array([      nan, 0.70909091,      nan, 0.70909091,
nan,
    0.67272727,      nan, 0.70909091]),
'mean_test_score': array([      nan, 0.80064935,      nan, 0.80428571,
nan,
    0.80422078,      nan, 0.80792208]),
'std_test_score': array([      nan, 0.05417001,      nan, 0.05925779,
nan,
    0.07425631,      nan, 0.06054289]),
'rank_test_score': array([5, 4, 6, 2, 7, 3, 8, 1], dtype=int32)}

```

```
[26]: grid_classifier_recall.cv_results_['mean_test_score']
```

```
[26]: array([      nan, 0.80064935,      nan, 0.80428571,      nan,
    0.80422078,      nan, 0.80792208])
```

```
[28]: np.array([grid_classifier_recall.cv_results_['mean_test_score'][x:x+2] for x in
    ↪range(0, len(grid_classifier_recall.cv_results_['mean_test_score']), 2)])
```

```
[28]: array([[      nan, 0.80064935],
    [      nan, 0.80428571],
    [      nan, 0.80422078],
    [      nan, 0.80792208]])
```

```
[29]: # from sklearn.model_selection import GridSearchCV
# from sklearn.linear_model import LogisticRegression

# lr = LogisticRegression().fit(X_train, y_train)
# grid_values = {'penalty': ['l1', 'l2'], 'C': [0.01, 0.1, 1, 10, 100]}
# grid_clf_rec = GridSearchCV(lr, param_grid = grid_values, scoring = 'recall')
# grid_clf_rec.fit(X_train, y_train)


# np.array([grid_clf_rec.cv_results_['mean_test_score'][x:x+2] for x in
    ↪range(0, len(grid_clf_rec.cv_results_['mean_test_score']), 2)])
```

```
[32]: # Use the following function to help visualize results from the grid search
def GridSearch_Heatmap(scores):
    %matplotlib widget
    import seaborn as sns
    import matplotlib.pyplot as plt
    plt.figure()
```

```

sns.heatmap(scores.reshape(4,2), xticklabels=['l1','l2'], yticklabels=[0.
↪01, 0.1, 1, 10])
plt.yticks(rotation=0);
GridSearch_Hemap(answer_six())

```



[]: