

Module 1

September 25, 2023

You are currently looking at **version 1.0** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](#) course resource.

0.1 Applied Machine Learning, Module 1: A simple classification task

0.1.1 Import required modules and load data file

```
[1]: # %matplotlib notebook
%matplotlib widget
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split

fruits = pd.read_table('assets/fruit_data_with_colors.txt')
```

```
[2]: fruits.head()
```

```
[2]:   fruit_label  fruit_name  fruit_subtype  mass  width  height  color_score
0           1      apple  granny_smith    192   8.4    7.3         0.55
1           1      apple  granny_smith    180   8.0    6.8         0.59
2           1      apple  granny_smith    176   7.4    7.2         0.60
3           2    mandarin      mandarin     86   6.2    4.7         0.80
4           2    mandarin      mandarin     84   6.0    4.6         0.79
```

```
[3]: # create a mapping from fruit label value to fruit name to make results easier
    ↪to interpret
lookup_fruit_name = dict(zip(fruits.fruit_label.unique(), fruits.fruit_name.
    ↪unique()))
lookup_fruit_name
```

```
[3]: {1: 'apple', 2: 'mandarin', 3: 'orange', 4: 'lemon'}
```

```
[17]: lookup_fruit_name[1]
```

```
[17]: 'apple'
```

The file contains the mass, height, and width of a selection of oranges, lemons and apples. The heights were measured along the core of the fruit. The widths were the widest width perpendicular to the height.

0.1.2 Examining the data

```
[4]: pd.__version__
```

```
[4]: '1.5.2'
```

```
[7]: # !jupyter lab --version
```

```
[8]: # !pip install --upgrade jupyterlab
```

```
[9]: # !ipython --version
```

```
[10]: # !pip install --upgrade ipython
```

```
[11]: # from IPython.display import display, Javascript
# display(Javascript("window.IPython = window.IPython || {}"))
```

```
[12]: # def execute_javascript(js_code):
#     from IPython.display import display, Javascript
#
#     wrapped_code = f"""
#     (function() {{
#         window.IPython = window.IPython || {{}};
#         {js_code}
#     }})();
#     """
#
#     display(Javascript(wrapped_code))
```

```
[13]: # execute_javascript("console.log('Hello, world!')")
```

```
[14]: # def is_running_in_ipython():
#     try:
#         __IPYTHON__
#         return True
#     except NameError:
#         return False
```

```
[15]: # is_running_in_ipython()
```

```

[16]: # !initializeIPython()

[17]: # !console.log(ipython)

[18]: # !typeof ipython

[19]: # !pip install ipympl

[5]: # %matplotlib widget

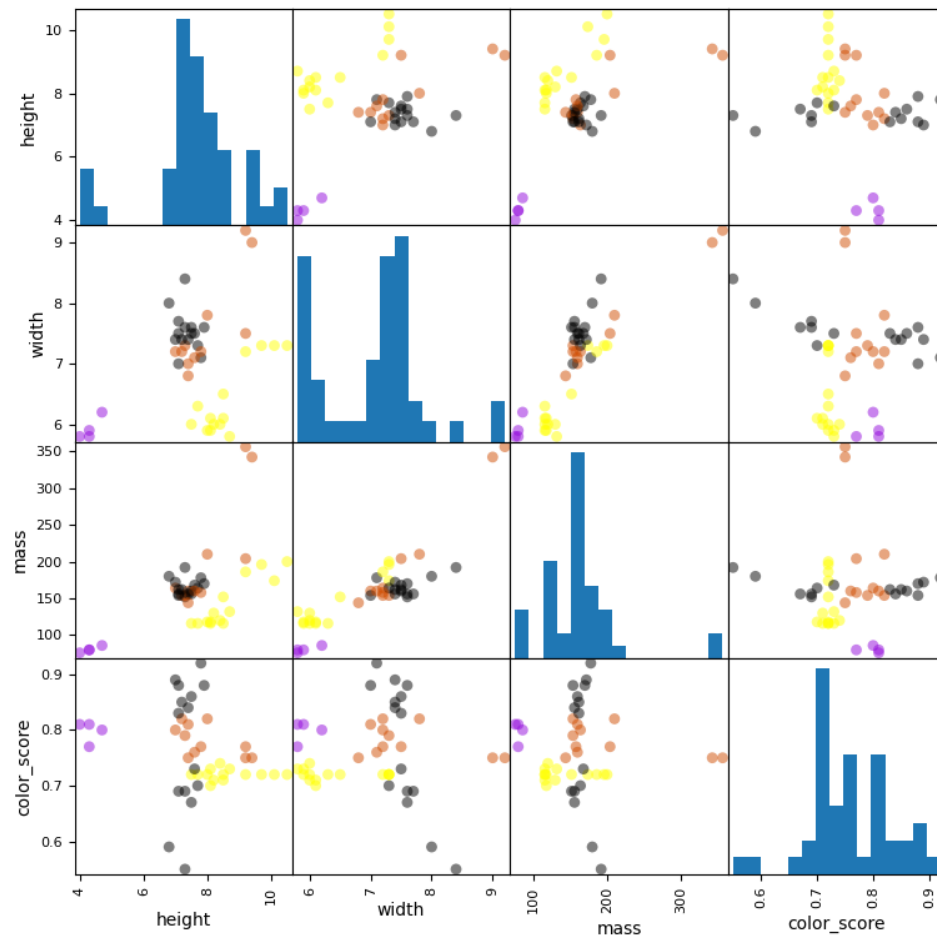
# this one solved the "JavaScript Error: IPython is not defined" error

[6]: # plotting a scatter matrix
from matplotlib import cm

X = fruits[['height', 'width', 'mass', 'color_score']]
y = fruits['fruit_label']
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)


cmap = cm.get_cmap('gnuplot')
scatter = pd.plotting.scatter_matrix(X_train, c= y_train, marker = 'o', s=40,
↪ hist_kws={'bins':15}, figsize=(9,9), cmap=cmap)

```



```
[7]: # plotting a 3D scatter plot
from mpl_toolkits.mplot3d import Axes3D

fig = plt.figure()
ax = fig.add_subplot(111, projection = '3d')
ax.scatter(X_train['width'], X_train['height'], X_train['color_score'], c = y_train, marker = 'o', s=100)
ax.set_xlabel('width')
ax.set_ylabel('height')
ax.set_zlabel('color_score')
plt.show()
```



0.1.3 Create train-test split

```
[8]: # For this example, we use the mass, width, and height features of each fruit_  
      ↳ instance  
X = fruits[['mass', 'width', 'height']]  
y = fruits['fruit_label']  
  
# default is 75% / 25% train-test split  
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

0.1.4 Create classifier object

```
[9]: from sklearn.neighbors import KNeighborsClassifier  
  
knn = KNeighborsClassifier(n_neighbors = 5)
```

0.1.5 Train the classifier (fit the estimator) using the training data

```
[10]: knn.fit(X_train, y_train)
```

```
[10]: KNeighborsClassifier()
```

0.1.6 Estimate the accuracy of the classifier on future data, using the test data

```
[11]: knn.score(X_test, y_test)
```

```
[11]: 0.5333333333333333
```

0.1.7 Use the trained k-NN classifier model to classify new, previously unseen objects

```
[20]: # first example: a small fruit with mass 20g, width 4.3 cm, height 5.5 cm
fruit_prediction = knn.predict([[20, 4.3, 5.5]])
lookup_fruit_name[fruit_prediction[0]]
```

```
[20]: 'mandarin'
```


```
[21]: # second example: a larger, elongated fruit with mass 100g, width 6.3 cm, ↵
      ↪height 8.5 cm
fruit_prediction = knn.predict([[100, 6.3, 8.5]])
lookup_fruit_name[fruit_prediction[0]]
```

```
[21]: 'lemon'
```

0.1.8 Plot the decision boundaries of the k-NN classifier

```
[22]: from adspy_shared_utilities import plot_fruit_knn

plot_fruit_knn(X_train, y_train, 5, 'uniform') # we choose 5 nearest neighbors
```




0.1.9 How sensitive is k-NN classification accuracy to the choice of the 'k' parameter?

```
[23]: k_range = range(1,20)
scores = []

for k in k_range:
    knn = KNeighborsClassifier(n_neighbors = k)
    knn.fit(X_train, y_train)
    scores.append(knn.score(X_test, y_test))

plt.figure()
plt.xlabel('k')
plt.ylabel('accuracy')
plt.scatter(k_range, scores)
plt.xticks([0,5,10,15,20]);
```



```
[27]: scores[:10]
```

```
[27]: [0.2916666666666667,
0.2708333333333333,
0.5208333333333334,
0.3958333333333333,
0.4375,
0.3333333333333333,
0.4791666666666667,
0.3541666666666667,
0.3958333333333333,
0.3541666666666667]
```

0.1.10 How sensitive is k-NN classification accuracy to the train/test split proportion?

```
[28]: t = [0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2]

knn = KNeighborsClassifier(n_neighbors = 5)

plt.figure()
```




```

for s in t:
    scores = []
    for i in range(1,1000):
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 1-s)
        knn.fit(X_train, y_train)
        scores.append(knn.score(X_test, y_test))
    plt.plot(s, np.mean(scores), 'bo')

plt.xlabel('Training set proportion (%)')
plt.ylabel('accuracy');

```



[]: